

# Project Introduction

- The primary goal is to process raw audio inputs and generate sentiment predictions through sequential speech-to-text transcription and sentiment classification.
- It aims to establish a modular and scalable architecture capable of integrating multiple ASR and sentiment models.
- The objective is to benchmark these model combinations based on accuracy, latency, throughput, and overall efficiency.
- Ultimately, it seeks to evolve into an operational and deployable system suitable for real-time audio sentiment analysis.

## Group Members

- Ziyang Jia
- Souvik Bag
- Srikanth Alla

## Settings

### Models Included

1. Audio to Text
  - AventIQ-AI/whisper-audio-to-text
    - Base Architecture: OpenAI Whisper
    - Dataset / Fine-Tuning: Mozilla Common Voice 13.0 dataset
    - Task / Use-Case: Speech-to-Text (Automatic Speech Recognition, ASR)
    - Parameter Size: 72M
1. Text to Sentiment
  - distilbert-base-uncased-finetuned-sst-2-english (Small)
    - Base Architecture: DistilBERT-base-uncased
    - Dataset / Fine-Tuning: Stanford Sentiment Treebank (SST-2)
    - Task / Use-Case: Sentiment Analysis
    - Parameter Size: 66M
  - siebert/sentiment-roberta-large-english (Large)
    - Base Architecture: RoBERTa-Large
    - Dataset / Fine-Tuning: 15 mixed English sentiment datasets
    - Task / Use-Case: Sentiment Analysis
    - Parameter Size: 355M
  - cardiffnlp/twitter-roberta-base-sentiment-latest (Medium)
    - Base Architecture: RoBERTa-Base
    - Dataset / Fine-Tuning: 124 M tweets (TweetEval benchmark)
    - Task / Use-Case: Sentiment Analysis
    - Parameter Size: 125M

## Model Variant

1. Quantitatively Evaluated
  - fp32 - audio to text models and text to sentiment models
  - fp16 - text to sentiment models
1. Provided in Web Service (will be quantitatively evaluated in Milestone 3)
  - fp32 - audio to text models and text to sentiment models
  - int8 - text to sentiment models

## Model Storage

- MinIO
  - An open-source, high-performance object storage server compatible with Amazon S3 APIs.
  - Usage in this pipeline: Store our models in ONNX

## Temporal File Storage

- Redis
  - A high-speed, in-memory key-value store that can act as a cache, database, or message broker.
  - Usage in this pipeline: Cache frequent inference results to reduce model load.

## Web Server

- FastAPI
  - A lightweight, asynchronous web framework built on Starlette (for async I/O) and Pydantic (for data validation).
  - Usage in this pipeline: deploy docker containers as web service.

## Flexible Container-Based Web Service:

- We kept both audio-to-text and text-to-sentiment models as flexible options when we create a container
- We assigned different ports to different model combinations
- Example:
  - input: .wav file
  - output:
    - transcribed text
    - sentiment classification
    - inference time

```

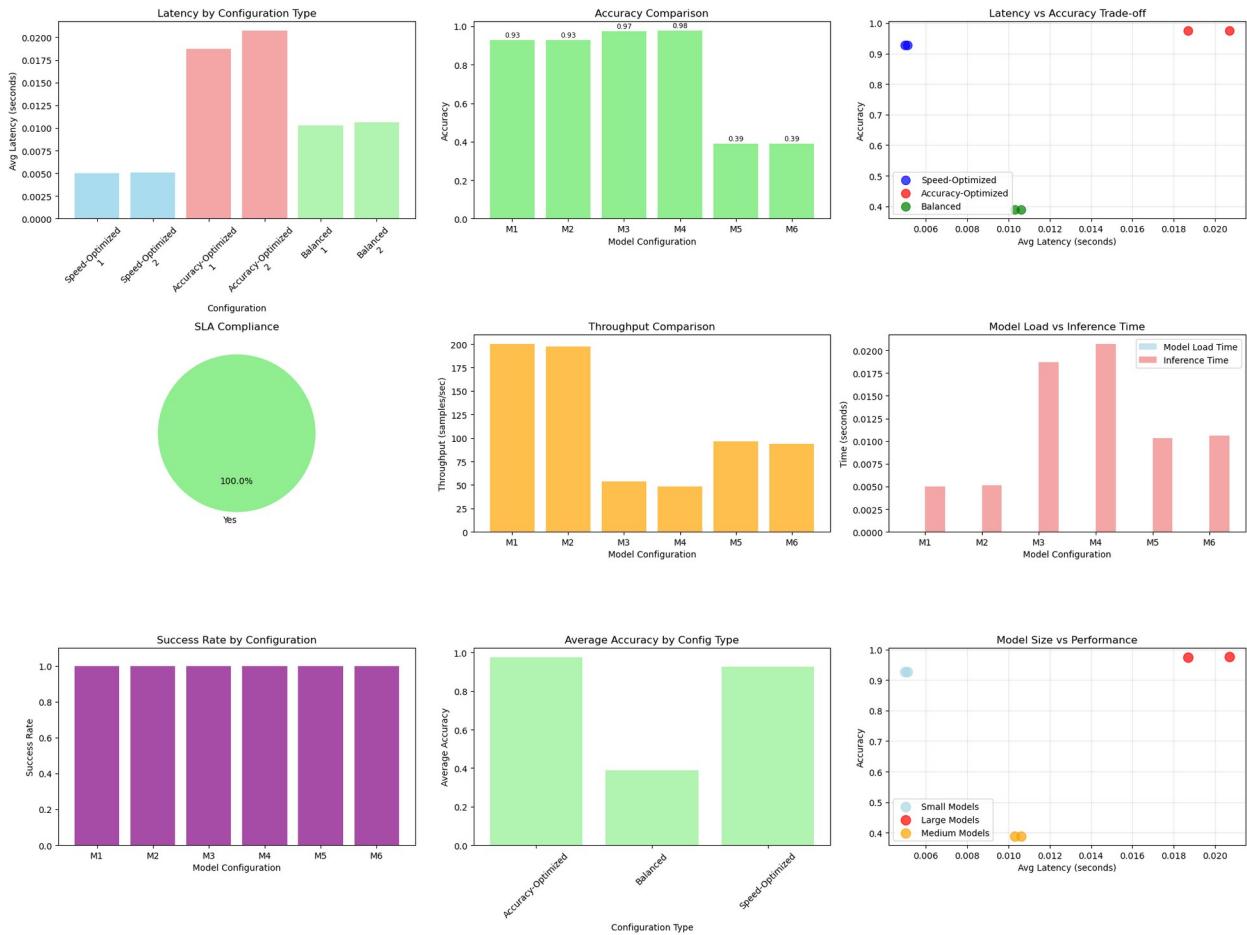
Windows PowerShell
+ - x
(C:\whisper_export) PS D:\Mizzou\25fall\eeml\project\e2eml-audio\ms_2> docker run -d -p 8002:8000
>> -e A2T_MODEL_PATH="/models/a2t_AventIQ-AI_whisper-audio-to-text_fp32"
>> -e T2S_MODEL_PATH="/models/t2s_distilbert-base-uncased-finetuned-sst-2-english_fp32.onnx"
>> -v ${PWD}:/app:/app
>> -v "${PWD}:/models:/models"
>> audio-sentiment
d9042b02240a979ed0b68e127cbf2a1ee6edb9b5bead1458387571cf1b962122
(C:\whisper_export) PS D:\Mizzou\25fall\eeml\project\e2eml-audio\ms_2> docker ps
CONTAINER ID IMAGE COMMAND CREATED STATUS PORTS
 NAMES
d9042b02240a979ed0b68e127cbf2a1ee6edb9b5bead1458387571cf1b962122
(C:\whisper_export) PS D:\Mizzou\25fall\eeml\project\e2eml-audio\ms_2> $FilePath = "D:\Mizzou\25fall\eeml\project\e2eml-audio\data\audio\audio_00000.wav"
(C:\whisper_export) PS D:\Mizzou\25fall\eeml\project\e2eml-audio\ms_2> $FileContent = [System.IO.File]::ReadAllBytes($FilePath)
(C:\whisper_export) PS D:\Mizzou\25fall\eeml\project\e2eml-audio\ms_2> $FileName = [System.IO.Path]::GetFileName($FilePath)
(C:\whisper_export) PS D:\Mizzou\25fall\eeml\project\e2eml-audio\ms_2> $boundary = "----WebKitFormBoundary$([System.Guid]::NewGuid().ToString())"
(C:\whisper_export) PS D:\Mizzou\25fall\eeml\project\e2eml-audio\ms_2>
(C:\whisper_export) PS D:\Mizzou\25fall\eeml\project\e2eml-audio\ms_2> $bodyLines = @(
>>    "--$boundary",
>>    "Content-Disposition: form-data; name='audio'; filename='$FileName''",
>>    "Content-Type: audio/wav",
>>    ""
>>    [System.Text.Encoding]::UTF8.GetString($FileContent),
>>    "--$boundary--"
>> ) -join "`r`n"
(C:\whisper_export) PS D:\Mizzou\25fall\eeml\project\e2eml-audio\ms_2>
(C:\whisper_export) PS D:\Mizzou\25fall\eeml\project\e2eml-audio\ms_2> Invoke-RestMethod -Uri "http://localhost:8002/predict"
>> -Method POST
>> -ContentType "multipart/form-data; boundary=$boundary"
>> -Body $bodyLines
transcript
sentiment
Merely, he is a man of the same name as the king of the world. @{label=POSITIVE; confidence=0.7689058184623718}

(C:\whisper_export) PS D:\Mizzou\25fall\eeml\project\e2eml-audio\ms_2> Measure-Command {
>>    $FilePath = "D:\Mizzou\25fall\eeml\project\e2eml-audio\data\audios\audio_00000.wav"
>>    $FileContent = [System.IO.File]::ReadAllBytes($FilePath)
>>    $FileName = [System.IO.Path]::GetFileName($FilePath)
>>
>>    $boundary = "----WebKitFormBoundary$([System.Guid]::NewGuid().ToString())"
>>
>>    $bodyLines = @(
>>        "--$boundary",
>>        "Content-Disposition: form-data; name='audio'; filename='$FileName''",
>>        "Content-Type: audio/wav",
>>        ""
>>        [System.Text.Encoding]::UTF8.GetString($FileContent),
>>        "--$boundary--"
>>    ) -join "`r`n"
>>
>>    Invoke-RestMethod -Uri "http://localhost:8002/predict"
>>    -Method POST
>>    -ContentType "multipart/form-data; boundary=$boundary"
>>    -Body $bodyLines
>> }

Days : 0
Hours : 0
Minutes : 0
Seconds : 0
Milliseconds : 821
Ticks : 8211151
TotalDays : 9.50364699074074E-06
TotalHours : 0.000228087527777778
TotalMinutes : 0.0136852516666667

```

# Evaluation Plots



In the plots above, we used the same audio-to-text model(AventIQ-AI/whisper-audio-to-text fp32) and different variants of text-to-sentiment models:

- M1: distilbert-base-uncased-finetuned-sst-2-english (fp32)
- M2: distilbert-base-uncased-finetuned-sst-2-english (fp16)
- M3: siebert/sentiment-roberta-large-english (fp32)
- M4: siebert/sentiment-roberta-large-english (fp16)
- M5: cardiffnlp/twitter-roberta-base-sentiment-latest (fp32)
- M6: cardiffnlp/twitter-roberta-base-sentiment-latest (fp16)

Variant	Scale	Avg Latency (s)	Accuracy	Throughput (samples/sec)	Relative Cost (~ Compute Load)	Notes
M1 (Speed-Optimized 1)	Small	0.005 s	0.93	200	\$ Low	Fastest; ideal for real-time use; small model size
M2 (Speed-Optimized 2)	Small	0.005 s	0.93	195	\$ Low	Similar to M1; nearly identical accuracy & latency
M3 (Accuracy-Optimized 1)	Large	0.018 s	0.97	55	\$ \$ \$ High	High-accuracy but slower; high compute cost
M4 (Accuracy-Optimized 2)	Large	0.021 s	0.98	50	\$ \$ \$ High	Highest accuracy; highest latency
M5 (Balanced 1)	Medium	0.010 s	0.39	100	\$ \$ Medium	Trade-off config; moderate speed, low accuracy
M6 (Balanced 2)	Medium	0.010 s	0.39	95	\$ \$ Medium	Similar to M5; balanced cost/speed but weak accuracy

## Empirical Conclusions:

- the differences on precision (fp32 vs fp16)
  - doesn't impact too much in terms of accuracy and throughput
  - impacts more on inference latency
- the model structure impacts more on temporal efficiency, e.g. given acc=0.93 and 0.97 for M1(fp16) and M3(fp16), we found that
  - the latency of M1 is only around 1/4 of M3
  - the throughput of M3 is only 1/4 of M1
  - this observation holds for other comparison pairs as well
- the model structure impacts more on accuracy as well, e.g. given parameter\_size=66M and 125M for M1(fp16) and M5(fp16), the temporal efficiency pattern still holds but we found:
  - the accuracy of M1 is more than twice of M5's accuracy
- some conclusions in model selection and optimization:
  - some model structure+pretrained dataset settings inherently get better accuracies in sentiment classification
  - precision impacts the accuracy but the difference is trivial in this task and in real-world applications, fp16 is more promising than the original fp32 models (in terms of temporal efficiency)

- inspirations for our future plan in milestone 3: we'll
  - push the temporal efficiency further to explore the capability of int8 models
  - try more flexible strategies for model choosing given different input (will consider some features like size of the input audio)
  - provide an interactive webpage if time permits